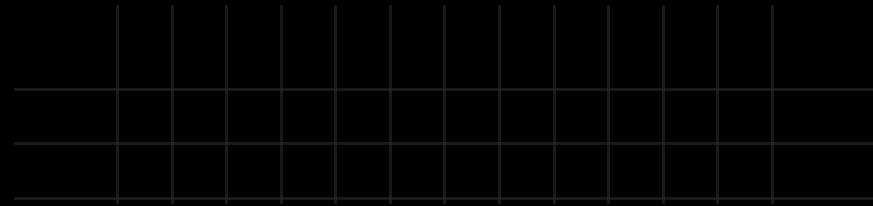


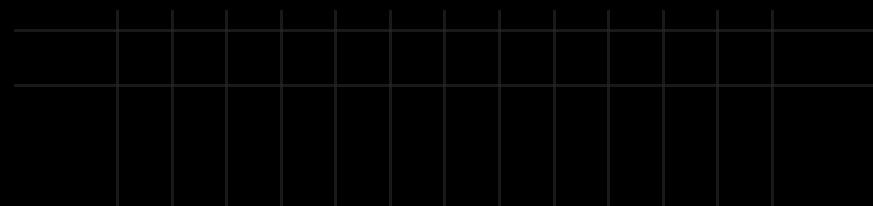
2025



DIGITAL IMAGE PROCESSING

2025 ASSIGNMENT 05

2023102070



INDIRA C REDDY
IIIT HYDERABAD

TABLE OF CONTENTS

- INTRODUCTION
- Q1 – READING IMAGE INTO ARRAY
- Q2 – WRITING ARRAY INTO IMAGE
- Q3 – BRIGHTNESS ADJUSTMENT
- Q4 – CONTRAST ADJUSTMENT
- Q5 – GRayscale CONVERSIONS
- Q6 – PSEUDOCOLOR MAPPING
- Q7 – GREEN SCREEN REPLACEMENT
- Q8 – VIDEO I/O
- Q9 – IMAGE TRANSITION VIDEO
- CONCLUSION & LEARNINGS

INTRODUCTION

This is a simple, first-pass tour of digital image processing using just **NumPy for math** and **Pillow/ImageIO for input/output**. No fancy libraries, just pixels and a few classic tricks.

The report is split into nine tiny tasks (Q1–Q9): read/write images, tweak brightness and contrast, convert RGB→grayscale (three ways), add a basic pseudocolor look, do a green-screen swap, flip between frames and videos, and build a clean fade transition. Everything stays in array form ($H \times W$ or $H \times W \times 3$) so it's easy to reason about what each line of code does. Goal was keep it clear, reproducible, and beginner-friendly. Each question shows the approach in one or two lines of math, the exact assumption (like value ranges or modes), and a quick observation of how the image changed.

[GitHub Repo link](#)

Q1

READING AN IMAGE INTO ARRAY

OBJECTIVE:

Read an image file and convert it into a NumPy array for processing.

APPROACH:

Use `PIL.Image.open()` to load file and convert to a NumPy array (`uint8`).

Supports grayscale ($H \times W$) and RGB ($H \times W \times 3$).

Print shape, type, and sample pixel values.

CODE SNIPPETS:

The screenshot shows a code editor window titled "q1.py". The code is as follows:

```
q1.py      x
Ubuntu-22.04 > home > indira > sem5_DIP_assignments > assign0 > questions > q1.py > ...
1  from utils.image_utils import read_image, save_pixel_array_to_txt
2
3  def run_q1():
4      path = "input_images/cat.jpeg"
5      img_array = read_image(path)
6
7      print("Image shape:", img_array.shape)
8      save_pixel_array_to_txt(img_array, "output_images/q1_output_pixels.txt")
9
#to read an image and return it as a numpy array.
#this supports grayscale and rgb too.
#q1
def read_image(path):
    img = Image.open(path)
    img_array = np.array(img)
    return img_array
#to save the pixel values of an image array in a .txt file.
def save_pixel_array_to_txt(array, output_path):
    with open(output_path, 'w') as f:
        if len(array.shape) == 2:
            #grayscale
            for row in array:
                f.write(' '.join(str(val) for val in row) + '\n')
        elif len(array.shape) == 3:
            #color
            for row in array:
                row_str = ' '.join(' '.join(str(channel) for channel in pixel) for pixel in row)
                f.write(row_str + '\n')
        else:
            raise ValueError("Unsupported array shape")
```

OUTPUT:

```
indira@IndiraCReddy:~/sem5_DIP_assignments/assign0$ python3 main.py
Image shape: (261, 193, 3)
```

Input and output images in next page.

OBSERVATIONS:

Image is stored as an array with dimensions (height, width, channels).

Image data is now matrix-form, easy to index and process.

The output array shown below is for a color image.

For a grayscale image each pixel is denoted with only one value while for a color image each pixel is shown with 3 nos which represent R, G and B



INPUT IMAGE USED

SAMPLE DATA RECEIVED

Q2

READING AN IMAGE INTO ARRAY

OBJECTIVE:

Convert a NumPy array back into an image and saving it

APPROACH:

Determine if array is grayscale or RGB based on its shape.

Use Pillow to create an image object from numpy array.

Save the output image to a specified output path.

CODE SNIPPETS:

```
# to write a numpy array into an image file
#again, supports grayscale(2D) and RGB color(3D) images too.
#q2
def write_image(array, output_path):
    if len(array.shape) == 2:
        mode = 'L' #grayscale
    elif len(array.shape) == 3 and array.shape[2] == 3:
        mode = 'RGB'
    else:
        raise ValueError("Unsupported image array shape")

img = Image.fromarray(np.uint8(a
img.save(output_path)
```

The screenshot shows a terminal window with several tabs at the top: main.py, image_utils.py 3, and q2.py 1 X. The current tab is q2.py. The code in the terminal is:

```
1  from utils.image_utils import read_image, write_image
2
3  def run_q2():
4      input_path = "input_images/cat.jpeg"
5      output_path = "output_images/q2_written_image.jpg"
6
7      img_array = read_image(input_path)
8      write_image(img_array, output_path)
9
10     print("Image successfully written to", output_path)
11
```

OUTPUT:

```
indira@IndiraCReddy:~/sem5_DIP_assignments/assign0$ python3 main.py
Image successfully written to output_images/q2_written_image.jpg
```

input and output images in next page.

OBSERVATIONS:

Values outside [0,255] would clip after casting. Output matches input layout (grayscale vs color).



INPUT IMAGE USED



OUTPUT - WRITTEN IMAGE

Q3

BRIGHTNESS ADJUSTMENT

OBJECTIVE:

Increase or decrease image brightness

APPROACH:

Load image into NumPy array.

For brightness: add a constant to all the pixel values

Clip values to [0, 255] before saving.

If 'value' added is positive we get a brighter image, if negative then we get a darker image.

Used int16 to avoid overflow and finalized as uint8

CODE SNIPPETS:

```
1 from utils.image_utils import read_image, write_image, change_brightness
2
3 def run_q3():
4     input_path = "input_images/cat.jpeg"
5     output_brighter = "output_images/q3_brighter.jpg"
6     output_darker = "output_images/q3_darker.jpg"
7
8     img = read_image(input_path)
9
10    brighter = change_brightness(img, 50) #increase brightness
11    darker = change_brightness(img, -50) #decrease brightness
12
13    write_image(brighter, output_brighter)
14    write_image(darker, output_darker)
15
16    print("Saved bright and dark versions!")
17
```

#to adjust brightness of image
#'value' if positive, we get a brighter image.
#if the 'value' is negative we get darker image
#q3

```
def change_brightness(img_array, value):
    #casting to int16 to avoid overflow/underflow cases
    img_int = img_array.astype(np.int16)
    #adding value, then clip to valid range
    brightened = np.clip(img_int + value, 0, 255)
    return brightened.astype(np.uint8)
```

OUTPUT:

```
indira@IndiraCReddy:~/sem5_DIP_assignments/assign0$ python3 main.py
Saved bright and dark versions!
```

Input and output images in next page.

OBSERVATIONS:

Adding a value or subtracting a value shifts intensity uniformly.

Large positive/negative shifts can cause highlight/shadow clipping.

INPUT IMAGE USED



BRIGHTER IMAGE



DARKED IMAGE



Q4

CONTRAST ADJUSTMENT

OBJECTIVE:

Modify the contrast of an image.

APPROACH:

Load image into NumPy array.

Apply scaling relative to the mean pixel value (i.e; multiplying around a midpoint)

Clip values to [0, 255] before saving.

Used float32 to avoid overflow and finalized as uint8

CODE SNIPPETS:

```
from image_utils import read_image, write_image, change_contrast

def run_q4():
    input_path = "input_images/cat.jpeg"
    output_high = "output_images/q4_high_contrast.jpg"
    output_low = "output_images/q4_low_contrast.jpg"

    img = read_image(input_path)

    high_contrast = change_contrast(img, 1.5) # more punch
    low_contrast = change_contrast(img, 0.5) # flatter

    write_image(high_contrast, output_high)
    write_image(low_contrast, output_low)

    print("Saved high and low contrast images!")

#adjusting the contrast of image
#multiplying around a midpoint
#if factor>1 then more contrast, if 0<factor<1 then less contrast
#q4
def change_contrast(img_array, factor):
    img_float = img_array.astype(np.float32)
    contrasted = (img_float - 128) * factor + 128
    return np.clip(contrasted, 0, 255).astype(np.uint8)
```

OUTPUT:

```
indira@IndiraCReddy:~/sem5_DIP_assignments/assign0$ python3 main.py
Saved high and low contrast images!
```

Input and output images in next page.

OBSERVATIONS:

Increasing factor >1 boosts contrast, <1 reduces it.

Higher contrast expands differences but can clip extremes; lower contrast compresses tones for a flatter look.

INPUT IMAGE USED



HIGH CONTRAST IMAGE



LOW CONTRAST IMAGE



Q5

GRAYSCALE CONVERSIONS

OBJECTIVE:

Convert a color image to grayscale using different methods

APPROACH:

Average method: $(R + G + B) / 3$

Luminosity method: weighted sum $(0.299R + 0.587G + 0.114B)$

Lightness method: $(\max(R, G, B) + \min(R, G, B)) / 2$

CODE SNIPPETS:

```
#q5
def rgb_to_grayscale(img_array, method='luminosity'):
    if len(img_array.shape) != 3 or img_array.shape[2] != 3:
        raise ValueError("Input must be a color (RGB) image")

    img = img_array.astype(np.float32)
    R, G, B = img[:, :, 0], img[:, :, 1], img[:, :, 2]

    if method == 'average':
        gray = (R + G + B) / 3
    elif method == 'luminosity':
        gray = 0.299 * R + 0.587 * G + 0.114 * B
    elif method == 'lightness':
        gray = (np.max(img, axis=2) + np.min(img, axis=2)) / 2
    else:
        raise ValueError("Unknown grayscale method")

    return gray.astype(np.uint8)
```

```
1  from utils.image_utils import read_image, write_image, rgb_to_grayscale
2
3  def run_q5():
4      input_path = "input_images/cat.jpeg"
5
6      avg = rgb_to_grayscale(read_image(input_path), method='average')
7      lum = rgb_to_grayscale(read_image(input_path), method='luminosity')
8      light = rgb_to_grayscale(read_image(input_path), method='lightness')
9
10     write_image(avg, "output_images/q5_grayscale_average.jpg")
11     write_image(lum, "output_images/q5_grayscale_luminosity.jpg")
12     write_image(light, "output_images/q5_grayscale_lightness.jpg")
13
14     print("Saved all grayscale versions!")
15
```

OUTPUT:

```
indira@IndiraCReddy:~/sem5_DIP_assignments/assign0$ python3 main.py
Saved all grayscale versions!
```

Input and output images in next page.

OBSERVATIONS:

Average: Equal weight to all channels. Simple but flat-looking.

Luminosity: Weighted for human perception (green-heavy). Looks more balanced and realistic.gives better perceptual accuracy than average.

Lightness: Based on max and min. Gives a contrasty, stylized feel.

INPUT IMAGE USED



AVERAGE GRayscale



LUMINOSITY GRayscale



LIGHTNESS GRayscale



Q6

PSEUDOCOLOR MAPPING

OBJECTIVE:

Apply false coloring to grayscale images for better visual interpretation.

APPROACH:

Convert original image to grayscale images

Function maps grayscale intensities to colors using a two-part scheme

- Values 0-127 are transformed into a spectrum from blue to red.
- Values 128-255 are transformed into a spectrum from red to yellow.

CODE SNIPPETS:

```
1  from utils.image_utils import read_image, write_image, rgb_to_grayscale, grayscale_to_pseudo
2
3  def run_q6():
4      input_path = "input_images/cat.jpeg"
5
6      gray = rgb_to_grayscale(read_image(input_path)) # default = luminosity
7      pseudo = grayscale_to_pseudocolor(gray)
8
9      write_image(pseudo, "output_images/q6_pseudocolor.jpg")#converts a grayscale image to a pseudocolor using a basic colormap
10     print("Pseudocolor image saved!")#output is hopefully an RGB image
11
12
13  def grayscale_to_pseudocolor(gray_img):
14      if len(gray_img.shape) != 2:
15          raise ValueError("Input must be a grayscale image")
16
17      height, width = gray_img.shape
18      pseudocolor = np.zeros((height, width, 3), dtype=np.uint8)
19
20      for i in range(height):
21          for j in range(width):
22              val = gray_img[i, j]
23              if val < 128:
24                  #map 0 - 127: Blue to Red
25                  ratio = val / 128
26                  r = int(255 * ratio)
27                  g = 0
28                  b = int(255 * (1 - ratio))
29              else:
30                  #map 128-255: Red to Yellow
31                  ratio = (val - 128) / 127
32                  r = 255
33                  g = int(255 * ratio)
34                  b = 0
35              pseudocolor[i, j] = [r, g, b]
36
37      return pseudocolor
```

OUTPUT:

```
indira@IndiraCReddy:~/sem5_DIP_assignments/assign0$ python3 main.py
Pseudocolor image saved!
```

Input and output images in next page.

OBSERVATIONS:

Got a heatmap kind of look as an output image.

Pseudocolor can highlight intensity variations better than plain grayscale.

Color mapping makes small intensity differences more visible than plain gray; shows smooth two-stage color ramp.

INPUT IMAGE USED



PSEUDOCOLOR IMAGE



Q7

GREEN SCREEN REPLACEMENT

OBJECTIVE:

Replace green background in a foreground image with another background

APPROACH:

Create a green mask with $(g > r * \text{threshold}) \& (g > b * \text{threshold})$. Resize background to foreground size, copy BG wherever mask is true.
Ensure both images are same sized.

CODE SNIPPETS:

```
replacing green pixels in foreground with background pixels
q7
def replace_green_screen(fg_img, bg_img, threshold=1.3):
    #removing alpha channel if present
    #alpha channel
    if fg_img.shape[2] == 4:
        fg_img = fg_img[:, :, :3]
    if bg_img.shape[2] == 4:
        bg_img = bg_img[:, :, :3]

    if fg_img.shape != bg_img.shape:
        bg_img = Image.fromarray(bg_img)
        bg_img = bg_img.resize((fg_img.shape[1], fg_img.shape[0])) #resizing to match (W, H)
        bg_img = np.array(bg_img)

    #print(fg_img.shape)
    #print(bg_img.shape)
    #raise ValueError("Foreground and background must be same shape")

    result = fg_img.copy()

    #extracting RGB channels
    r, g, b = fg_img[:, :, 0], fg_img[:, :, 1], fg_img[:, :, 2]

    #creating a green mask
    green_mask = (g > r * threshold) & (g > b * threshold)

    #applying mask - replace green pixels with background
    result[green_mask] = bg_img[green_mask]

    return result

from utils.image_utils import read_image, write_image, replace_green_screen

def run_q7():
    fg_path = "input_images/foreground.png"
    bg_path = "input_images/background.jpg"

    fg = read_image(fg_path)
    bg = read_image(bg_path)

    output = replace_green_screen(fg, bg)
    write_image(output, "output_images/q7_minions_in_the_sky.jpg")

    print("Green screen replaced and saved!")
```

OUTPUT:

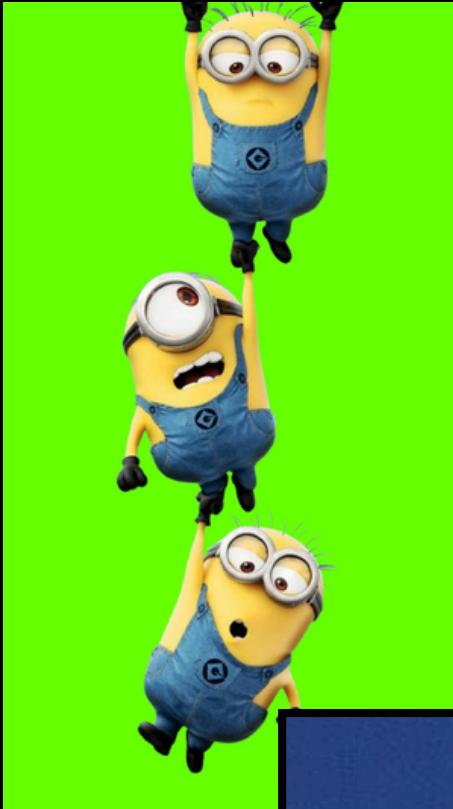
```
indira@IndiraCReddy:~/sem5_DIP_assignments/assign0$ python3 main.py
Green screen replaced and saved!
```

Input and output images in next page.

OBSERVATIONS:

Successfully replaced the green pixels with background pixels
Works well on clean green areas. May pick up green-ish clothes/objects;
fine-tune threshold based on scene.

GREEN SCREEN IMAGE



BACKGROUND IMAGE



+

↓↓



RESULT - MINIONS IN THE SKY

Q8

VIDEO I/O

OBJECTIVE:

Extract frames from video and reconstruct back.

APPROACH:

Read video using imageio

Save each image frame as .jpg

Recombine into .mp4 file using imageio.mimsave

CODE SNIPPETS:

```
1  from utils.image_utils import extract_frames_from_video, frames_to_video
2
3  def run_q8():
4      video_path = "videos/sample_video.mp4"
5      frame_folder = "output_images/q8_frames"
6      output_video = "videos/q8_output_video.mp4"
7
8      num_frames = extract_frames_from_video(video_path, frame_folder)
9      print(f"Extracted {num_frames} frames.")
10
11     frames_to_video(frame_folder, output_video)
12     print(f"Reconstructed video saved to {output_video}")
13 
```

```
#q8
#reading a video file and save each frame as an image into an output_fold
def extract_frames_from_video(video_path, output_folder):
    reader = imageio.get_reader(video_path)
    os.makedirs(output_folder, exist_ok=True)

    for i, frame in enumerate(reader):
        frame_path = os.path.join(output_folder, f"frame_{i:03d}.jpg")
        imageio.imwrite(frame_path, frame)

    return i + 1 #total number of frames
#reading all frames in input folder and then writing them into a video
#assuming all the frames are names in order like frame_000.jpg,etc.
def frames_to_video(input_folder, output_video_path, fps=24):
    writer = imageio.get_writer(output_video_path, fps=fps)

    frame_files = sorted(
        [f for f in os.listdir(input_folder) if f.endswith('.jpg')])
    )

    for filename in frame_files:
        frame = imageio.imread(os.path.join(input_folder, filename))
        writer.append_data(frame)

    writer.close()
```

OUTPUT:

```
indira@IndiraCRReddy:~/sem5_DIP_assignments/assign0$ python3 main.py
Extracted 68 frames.
```

Input video link and output image & video links in next page.

OBSERVATIONS:

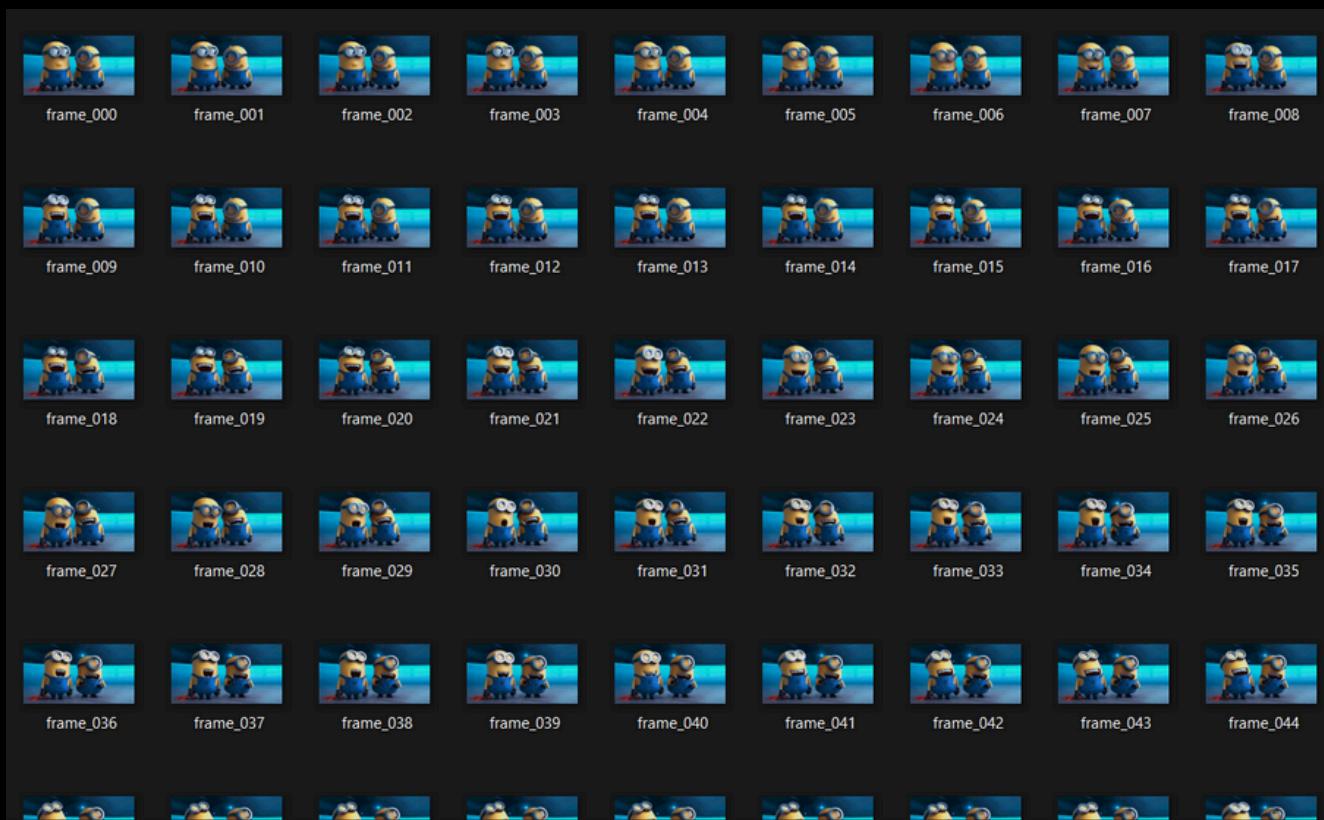
Video was segmented into 68 frames and successful in reconstructing the video from said frames.

Frames are assumed sequentially named and same size; FPS controls output duration.

Must ensure frame dimensions are divisible by 16 for best codec compatibility.

INPUT VIDEO LINK

SAMPLE FRAMES OBTAINED



OUTPUT VIDEO LINK

Q9

IMAGE TRANSITION VIDEO

OBJECTIVE:

Create a fade transition between two images

APPROACH:

Load two images of same size (resize second image to first if needed)
Generate num_frames blends with alpha = i/(num_frames-1) and blend = (1- α)*img1 + α *img2.
Blend them progressively over those frames
Save as .mp4 file using imageio.

CODE SNIPPETS:

```
#q9
#creates a list of frames fading from img1 to img2
def create_fade_transition(img1, img2, num_frames=24):
    if img1.shape != img2.shape:
        img2 = Image.fromarray(img2)
        img2 = img2.resize((img1.shape[1], img1.shape[0]))
        img2 = np.array(img2)
    if img1.shape != img2.shape:
        raise ValueError("Images must be the same shape for transition")

    img1 = img1.astype(np.float32)
    img2 = img2.astype(np.float32)
    frames = []

    for i in range(num_frames):
        alpha = i / (num_frames - 1)
        blended = (1 - alpha) * img1 + alpha * img2
        frames.append(blended.astype(np.uint8))

    return frames
```

```
from utils.image_utils import read_image, create_fade_transition
import imageio
import numpy as np
from PIL import Image

def run_q9():
    img1 = read_image("input_images/transition_start.jpg")
    img2 = read_image("input_images/transition_end.jpg")

    # Resize img2 to match img1 if needed
    if img1.shape != img2.shape:

        img2 = Image.fromarray(img2).resize((img1.shape[1], img1.shape[0]))
        img2 = np.array(img2)

    frames = create_fade_transition(img1, img2, num_frames=72) #3 sec at 24fps
    output_path = "videos/q9_transition_video.mp4"
    writer = imageio.get_writer(output_path, fps=24)

    for frame in frames:
        writer.append_data(frame)
    writer.close()

    print(f"Transition video saved to {output_path}")
```

OUTPUT:

```
indira@IndiraCReddy:~/sem5_DIP_assignments/assign0$ python3 main.py
IMAGEIO FFMPEG_WRITER WARNING: input image is not divisible by macro_block_size=16, resizing from (613, 407) to (624, 416) to ensure video compatibility with most codecs and players. To prevent resizing, make your input image divisible by the macro_block_size or set the macro_block_size to 1 (risking incompatibility).
Transition video saved to videos/q9_transition_video.mp4
```

Input images and output video link in next page.

OBSERVATIONS:

Produces a smooth crossfade. More frames = smoother transition;
mismatched sizes were handled by resizing.
Duration depends on frame count and FPS.

INPUT IMG1 - TRANSITION START



INPUT IMG2 - TRANSITION END



[OUTPUT VIDEO LINK](#)

CONCLUSION

Mission accomplished: treated images as plain NumPy arrays, did safe brightness/contrast tweaks, compared three grayscale styles (luminosity looked the most natural), added a simple two-step pseudocolor to reveal subtle differences, replaced green backgrounds with a clean RGB-ratio mask, and handled frames \leftrightarrow video plus a smooth crossfade.

It's intentionally basic, so a few tradeoffs are expected (looped pseudocolor, simple green key, fixed-size/FPS assumptions). For the next round, I'd keep the same structure and add tiny quality-of-life upgrades like before/after histograms and a one-liner CLI for reproducibility.

Key Learnings

- Arrays > magic: treating images as $H \times W$ / $H \times W \times 3$ makes every effect feel obvious and debuggable.
- Dtypes matter: do math in a wider type (int16/float32), then clip and cast to uint8 - no overflow drama.
- Clipping is real: big brightness/contrast pushes can crush shadows/highlights; “how much is too much” shows up in histograms.
- Grayscale isn't one thing: Average is simple but flat; Luminosity matches human perception better; Lightness can look stylized.
- Color maps reveal detail: even a basic blue \rightarrow red \rightarrow yellow ramp makes subtle intensity changes pop more than plain gray.
- Masking is a heuristic: the green-screen ratio works on clean screens; threshold tuning is normal (and scene-dependent).
- Video is just time + order: consistent naming, consistent size, and a chosen FPS fully define the output feel.
- Small structure = big clarity: one folder per task, predictable I/O paths, and tiny functions make the pipeline easy to test and explain.